

OpenMP Examples - Tasking

Marco Ferretti, Mirto Musci

Dipartimento di Ingegneria Industriale e dell'Informazione
University of Pavia

December 3, 2018

Outline

1 Tasking

- Basics
- Advanced

2 Exercises

- Assignment 2: Quicksort
- Assignment 3: Jacobi

Outline

1 Tasking

- Basics
- Advanced

2 Exercises

- Assignment 2: Quicksort
- Assignment 3: Jacobi

Why task parallelism?

List traversal

```
void traverse_list (List l)
{
    Element e ;
    #pragma omp parallel private(e)
    {
        for ( e=l->first; e; e=e->next )
            #pragma omp single nowait
            process(e);
    }
}
```

Without tasks

- Ackward
- Very poor performance
- Not composable

Why task parallelism?

Tree traversal

```
void traverse ( Tree *tree )
{
    #pragma omp parallel sections
    {
        #pragma omp section
        if ( tree->left )
            traverse ( tree->left );

        #pragma omp section
        if ( tree->right )
            traverse ( tree->right );
    }
    process ( tree );
}
```

Without tasks

- Too many parallel regions
- Extra overheads
- Extra synchronizations
- Not always well supported

Task parallelism

- Better solution for those problems
- Tasks were first introduced to OpenMP in version 3.0
- In general: **parallelize irregular problems**
 - unbounded loops
 - recursive algorithms
 - producer/consumer schemes
- In OpenMP 4.0, the `depend` clause and the `taskgroup` construct were introduced
 - `taskloops` were introduced in OpenMP 4.5

List traversal (with task)

Example

```
void traverse_list (List l)
{
    Element e;
    for (e=l->first; e; e=e->next)
        #pragma omp task
        process(e);
}
```

What is a task?

- Tasks are **code blocks** that the compiler wraps up and makes available to be executed in parallel
 - A task is a **package** of both code and variable
- Execution is completely in the hand of the compiler
 - No need to manually assign to threads!
 - Highly composable. Can be nested inside parallel regions, other tasks and worksharings

```
#pragma omp parallel
{
    #pragma omp task
    printf("hello world from a random thread\n");
}
```


How to create them?

- Like worksharing constructs, tasks must also be created inside of a parallel region.
- In the previous example **every thread** created a “print” task.
- To only spawn a task **once**, the **single** construct is used.
 - The execution is then **distributed to every thread**

```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task
        printf("hello world\n");

        #pragma omp task
        printf("hello again!\n");
    }
}
```

Task synchronization

- The previous example will only print *hello world* once, but *the ordering of hello world and hello again is undefined*.
 - Only one guarantee: both tasks will end at the parallel region barrier (or any other barrier)
- Two ways to specify the order: *taskwait* and *dependencies*
 - *taskwait* waits for tasks spawned by current task and itself

```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task
        printf("hello world\n");

        #pragma omp taskwait

        #pragma omp task
        printf("hello again!\n");
    }
}
```

Data sharing for tasks

- Like other OpenMP constructs, variables that are used inside a task can be specified:
 - **explicitly** with clauses (shared, firstprivate, private, etc.)
 - **implicitly**, if not otherwise specified.
- In general, data accessed by a task is **shared**.

Example 1

```
#pragma omp parallel
{
    int x = 0;
    #pragma omp single
    {
        #pragma omp task
        {
            x++;
            printf("from task 1: x = %d\n", x); }

        #pragma omp taskwait

        #pragma omp task
        {
            x++;
            printf("from task 2: x = %d\n", x); }
    }
}
```

- Variable `x` will be 1 and then 2, since both tasks are using the (implicitly) shared variable `x`.

Example 1

```
#pragma omp parallel
{
    int x = 0;
    #pragma omp single
    {
        #pragma omp task
        {
            x++;
            printf("from task 1: x = %d\n", x);
        }

        #pragma omp taskwait

        #pragma omp task
        {
            x++;
            printf("from task 2: x = %d\n", x);
        }
    }
}
```

- Variable `x` will be 1 and then 2, since both tasks are using the (implicitly) shared variable `x`.

Example 2

```
#pragma omp parallel
{ int x = 0;
  #pragma omp single
  {
    #pragma omp task firstprivate(x)
    { x++;
      printf("from task 1: x = %d\n", x); 10 }

    #pragma omp taskwait

    #pragma omp task firstprivate(x)
    { x++;
      printf("from task 2: x = %d\n", x); }
  }
}
```

- Variable `x` is made `firstprivate`, so both tasks increment their copy of `x`, and both print `x=1`
 - `private` wouldn't be valid: `x` is not initialized in the task

Example 2

```
#pragma omp parallel
{ int x = 0;
  #pragma omp single
  {
    #pragma omp task firstprivate(x)
    { x++;
      printf("from task 1: x = %d\n", x); 10 }

    #pragma omp taskwait

    #pragma omp task firstprivate(x)
    { x++;
      printf("from task 2: x = %d\n", x); }
  }
}
```

- Variable `x` is made `firstprivate`, so both tasks increment their copy of `x`, and **both print `x=1`**
 - `private` wouldn't be valid: `x` is **not initialized** in the task

Advanced data sharing

```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task
        {
            int x = 0;
            #pragma omp task
            { x++;
              printf("x = %d\n"); }
        }
    }
}
```

- One task spawning another.
 - the **inner** task accesses a variable local to the **outer** task.
 - even if only one task modify it, x exists on the **stack of the outer task**: outer task can finish before inner task begins!
 - up to the programmer to synchronize (e.g. taskwait)

Outline

1 Tasking

- Basics
- **Advanced**

2 Exercises

- Assignment 2: Quicksort
- Assignment 3: Jacobi

The depend clause

- The depend clause takes a **type** followed by a **list of variables**

```
#pragma omp task depend(in: x)  
    depend(out: y) depend(inout: z)
```

- They are both to **correctly order the tasks**.
 - **IN** will make a task dependent on the last task that used the same variable as an out (or inout) dependency.
 - **OUT** will make a task dependent on the last task that used the same variable as an in, out (or inout) dependency
 - **INOUT** the same as OUT, only used for readability.
- There is no data movement with respect to external accesses

Example: simple depend

```
#pragma omp parallel
{
    #pragma omp single
    {
        int x, y, z;
        #pragma omp task depend(out: x)
        x = init();

        #pragma omp task depend(in: x) depend(out: y)
        y = f(x);

        #pragma omp task depend(in: x) depend(out: z)
        z = g(x);

        #pragma omp task depend(in: y, z)
        finalize(y, z);
    }
}
```

Example: array sections as dependencies

```
void matmul_depend(int N, int BS, float A[N][N], float B[N][N], float C[N][N])
{
    for(int i=0; i<N; i+=BS)
        for(int j=0; j<N; j+=BS)
            for(int k=0; k<N; k+=BS)

                #pragma omp task depend(in: A[i:BS][k:BS], B[k:BS][j:BS]) \
                depend(inout: C[i:BS][j:BS])

                for(int ii=i; ii<i+BS; ii++)
                    for(int jj=j; jj<j+BS; jj++)
                        for(int kk=k; kk<k+BS; kk++)
                            C[ii][jj] += A[ii][kk] * B[kk][jj];
}
```

- Assumptions:

- BS divides evenly into N
- i, j, k, A, B, C are firstprivate by default
- A, B and C are just pointers: they all refer to the same data.

Taskgroup I

- A taskgroup is similar to taskwait
 - waits on all descendant created in the block that follows
 - not only for childs of current task, like taskwait

```
#pragma omp taskgroup  
#pragma omp task  
    task_spawning_function();
```

Taskgroup II

- Taskgroups can also be used to selectively synchronize tasks:
 - background can continue to the end of the parallel region
 - task_spawning and all of its descendants will be waited at the end of the taskgroup.

```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task
        background_work();
        #pragma omp taskgroup
        {
            #pragma omp task
            task_spawning_function();
        }
    }
}
```

The taskloop construct

- `taskloop` precedes a `for` loop, and creates a tasks for one or more iterations of the loop.
 - `implicit` taskgroup
 - basic usage is similar to a `for` construct
 - clauses from both `for` and `task` constructs for fine-tuning
- Unique clauses
 - `grainsize` and `num_tasks` are used to specify how to break up the iterations (similar to `chunksize` and number of threads)
 - `nogroup`: similar to `nowait`, removing the taskgroup behavior

```
#pragma omp taskloop
for(int i=0; i < N; i++)
    //Work
```

Outline

1 Tasking

- Basics
- Advanced

2 Exercises

- Assignment 2: Quicksort
- Assignment 3: Jacobi

Quicksort Algorithm

Algorithm

- Given an array of n elements (e.g., integers).
- If array only contains one element: **return**.
- Else: Pick one element to use as **pivot**.
- **Partition** elements into **two sub-arrays**:
 - Elements less than or equal to pivot
 - Elements greater than pivot
- **Recursively** quicksort two sub-arrays
- Return results

Notes

- There are a number of ways to pick the pivot element
 - Commonly first or last element, but bad performance if the array is already ordered
 - Random index or middle-point index solve the problem
- After partitioning, the sub-arrays can be **stored in the original data array** to increase memory efficiency
 - Partitioning loops through, swapping elements

Example

9 3 4 220 1 3 10 5 8

Choose a pivot.

9 3 4 220 1 3 10 5 8

Partition data by pivot value.

3 4 1 3 5 8 9 220 10

Sort each partitioned set.

⋮

1 3 3 4 5 8 9 10 220

Serial Code

```
void quicksort ( int a[ ], int lower, int upper )
{
    int i ;
    if ( upper > lower )
    {
        i = partition ( a, lower, upper ) ;
        quicksort ( a, lower, i - 1 ) ;
        quicksort ( a, i + 1, upper ) ;
    }
}
```

Assignment

- Refine the serial implementation provided
- Try to parallelize the code using OpenMP... is not easy!
 - Try with `section` constructs, or experiment with `task`
- Remember the code is `recursive`!
 - You could try to call `omp_set_nested(1)`
 - Or somehow limit thread spawning
- Also could be necessary to reduce task spawning... `if` clause (check documentation)
- Always measure performance with `omp_get_wtime`

Outline

1 Tasking

- Basics
- Advanced

2 Exercises

- Assignment 2: Quicksort
- Assignment 3: Jacobi

The long way to the final project...

- This final exercise will not contain any help, **just the source code** of a serial implementation
 - Has a level of complexity on par of a **final project**
- In the **final report** you should:
 - analyze the serial algorithm
 - implement it serially
 - analyze all of the possible parallel implementations
 - actually implement one or more of them
 - correctly measure the performances scaling the number of threads, the volume of data...

Source code

In numerical linear algebra, the Jacobi method is an algorithm for determining the solutions of a diagonally dominant system of linear equations.

- Here we only report the **core** of the serial algorithm
- As typical in parallelism, we **don't need to understand the meaning**; just the structure

```
for(int i=0; i<size; i++) {  
    for(int j=i+1; j<size; j++) {  
        A[j*size+i] /= A[i*size+i];  
        for(int k=i+1; k<size; k++) {  
            A[j*size+k] -= A[j*size+i] * A[i*size+k];  
        }  
    }  
}
```


Exercise

- Measure serial performance.
 - How can you parallelize it?
 - Be careful of dependencies.
- Try using the `for construct`... but where?
- Try using tasking and the `depend clause`
 - **Hint**: you have to protect three IN values and one OUT value
 - Strive for a `small granularity` with task!
- There are solutions provided... but **do not** look at them!
 - The second task solution is useful for huge dataset.. why?
 - **Hint**: it's about granularity...

For Further Reading



Blaise Barney

OpenMP Exercise, 2015

<https://computing.llnl.gov/tutorials/openMP/exercise.html>



NERCS

OpenMP Tasking Tutorial

<http://www.nersc.gov/users/software/programming-models/openmp/openmp-tasking/>